

Type-Aware Failure Prediction for Chart and Diagram Reasoning in Vision-Language Models: Mid-Term Report

Yash More, Naman Singhal, Shrish Kadam, Sanjith Ganapathi

{yash.more, naman.s, shrish.kadam}@research.iiit.ac.in, sanjith.ganapathi@students.iiit.ac.in
International Institute of Information Technology, Hyderabad

Abstract

When a vision-language model (VLM) answers a chart question wrongly, it may have misread the figure (a *structural* failure) or stated something the figure does not contain (a *fabrication*). We ask whether the two are distinguishable in the model’s internal state before it generates an answer. We ran Qwen2.5-VL-7B and LLaVA-NeXT-7B over the 2,500 ChartQA test questions and cached pooled activations at every decoder layer from the same forward pass that produced each answer. A linear probe on the last prompt token predicts whether the answer will be wrong with test AUROC 0.932 for Qwen (0.786 for a surface-feature baseline) and 0.798 for LLaVA (0.608). Labelling Qwen’s 317 errors showed that fabrication is essentially absent from ChartQA (1 of 317), and that 35% of the scored errors are not errors at all. We therefore built a synthetic bar-chart arm that asks about categories absent from the chart. Both models almost never refuse such questions (15 refusals in 1,200 for Qwen, 0 for LLaVA), yet both represent that the asked category is missing (probe AUROC 0.96 to 1.00, with text-only and image-only controls at chance). In Qwen, probes for misreads and for missing-bar fabrications are largely type-specific (own-type AUROC 0.887 and 1.000; cross-type 0.594 and 0.576). In LLaVA the missing-bar signal is specific but the misread probe is not: it fires more on fabrications (0.879) than on misreads (0.776), and a one-number chart feature predicts LLaVA’s misreads better than its activations. The Type Separability Hypothesis is thus supported for one model and only partly for the other.

1 Introduction

VLMs are increasingly used to read charts, plots and diagrams whose answers feed decisions. When such a model answers incorrectly, it can be wrong in two mechanically different ways. It may have **misread the figure**: picked the wrong bar, misjudged an axis tick, traced the wrong edge. The

correct answer was in the image, and a more careful look might recover it. Or it may have **fabricated**: stated a value or label that appears nowhere in the figure. Looking again will not help, and the answer should be withheld.

Reliability tooling today collapses both into one “do not trust” verdict. Abstaining on everything discards recoverable answers; re-querying everything spends compute on answers that were never grounded. Our project asks whether the distinction is visible in the model’s internal state *before* generation, and whether knowing it is worth anything.

Refined problem statement. Let $x = (I, q)$ be a chart image and a question, and $\hat{y}$ the answer a frozen VLM generates greedily. Let $h_\ell^p(x) \in \mathbb{R}^d$ be the model’s hidden state after decoder layer ℓ , pooled at prompt position p , computed in the prefill pass that precedes the first answer token. Each wrong answer carries a failure type $t \in \{\text{STRUCTURAL}, \text{FABRICATION}\}$ determined from $\hat{y}$ and the figure. We test (i) whether $\Pr(\hat{y} \text{ wrong} \mid h)$ and $\Pr(t \mid h, \hat{y} \text{ wrong})$ can be estimated by linear probes on h_ℓ^p better than by surface features of the input, and (ii) whether the probe for one failure type also responds to the other (cross-transfer). Later stages (E4, E5) test whether the type determines which repair works.

Progress summary. All planned inference and activation caching is complete for both models on ChartQA and on two synthetic pilots (Section 4). The binary and type probes (E1, E3) are trained with surface baselines (a partial E2). Section 5 reports the naturalistic arm, including why it cannot supply fabrication examples; Section 6 the synthetic arm built in response; Section 7 what the results say about the hypothesis; and Section 8 the plan for the remaining weeks.

2 Background and Related Work

2.1 Pre-generation failure detection

Kogilathota et al. (2026) establish that hallucination risk in VLMs is readable from internal representations in a single forward pass, without decoding, reaching up to 0.93 AUROC across eight models. Two of their findings shaped our design: the most informative representation family varies by architecture, with Qwen2.5-VL-7B best served by visual-only features at roughly 0.79 AUROC, where most models rely on late query-token states. Zhang et al. (2026) localise the signal further, isolating hallucination-relevant attention heads through a variational information bottleneck and using probe gradients to intervene at inference time. Both works address binary object hallucination on natural photographs. Neither asks whether distinct kinds of failure are internally distinguishable.

2.2 Evidence that failure types differ mechanistically

Liu et al. (2026) apply activation patching across five VLMs and identify a visual grounding pathway supporting correct predictions alongside a separate hallucination pathway driving erroneous ones, the former at mid-to-late depths and the latter concentrated at early layers and network boundaries. The identified circuit transfers selectively to relational but not attribute hallucination. That selectivity is the closest existing evidence that hallucination is not one internal phenomenon, and is the empirical basis for our hypothesis. It also motivates probing across the full depth of the network, since the two pathways do not occupy the same layers.

2.3 Chart and diagram reasoning

ChartQA (Masry et al., 2022) remains the standard benchmark for chart question answering with visual and logical reasoning. Rani et al. (2025) contribute chart, question, reasoning-step and attribution annotations, identifying which regions justify a model’s answer. Iyengar et al. (2026) assemble diagram QA instances across six source datasets, annotating the minimal regions required to verify each answer, and find that correct answers do not guarantee a model grounded its reasoning in the supporting regions. Both resources annotate *where a model should look*. Neither labels *how a model failed*, which is the supervision signal we need.

2.4 Error taxonomies and visual self-correction

Ashury-Tahan et al. (2026) argue that benchmark accuracy reports whether a model fails but not why, and build an LLM-judge pipeline for taxonomy construction across 35 datasets and 83 models. Their setting is text-only LLMs; we borrow the methodology, not the taxonomy. Li et al. (2026) propose Visual Self-Refine, in which a chart-parsing model emits pixel-level localisations, visualises them, and feeds them back to itself to correct its own perception errors. Making a model re-examine a figure is therefore established; what is untested is whether the *appropriate* repair depends on which failure occurred.

2.5 Probe validity

Two results define our standard of evidence. Sahoo et al. (2026) report probes separating reasoning types at 100% accuracy with well-separated geometry, then show the separation was driven entirely by format confounds, collapsing to chance once source identity, option count and response length were residualised out. Yuan et al. (2026) report the healthy contrast, a hidden-state probe at 0.95 AUROC against a text-surface classifier at 0.59, but find that four interventions built on that signal failed: the signal was diagnostic, not causal. Two lessons follow: a probe result is uninterpretable without a surface baseline, and detection does not imply a usable intervention.

2.6 Positioning

Pre-generation detection is binary and targets natural images; mechanistic type-selectivity has been shown for object hallucination, not structured-image reasoning; chart resources annotate evidence regions, not failure types; and visual self-correction is applied uniformly. Our work sits in the intersection: failure-*type* prediction on structured images, and whether that type should determine the repair.

3 Hypothesis

Type Separability Hypothesis. The pre-generation internal state preceding a structural misreading is linearly distinguishable from the state preceding a fabrication, because the former is a precision-limited perceptual failure while the latter reflects generation ungrounded in the image. We test this rather than assume it: a misread bar height may be a higher-precision instance of the same per-

ceptual failure, and the design treats a collapse of the two classes as a reportable outcome.

4 Experimental Setup

4.1 Models and inference

We use two frozen open VLMs: Qwen2.5-VL-7B (Bai et al., 2025) and LLaVA-NeXT with a Mistral-7B language model (llava-v1.6-mistral-7b; Liu et al., 2024). ChartGemma, the stretch goal, is dropped. Every question is asked with the standard ChartQA suffix “Answer the question using a single word or phrase.”, decoded greedily (at most 32 new tokens), in fp16 with SDPA attention on two RTX 2080 Ti GPUs. Hardware, library versions and settings are frozen and recorded with every run, because a different GPU kernel can change the generated answer and decouple the cached activations from the answer that was labelled. Answers are scored by ChartQA relaxed accuracy (numeric answers within 5% of the gold; strings by exact match).

4.2 Activation caching

Activations are captured by forward hooks *during the same* generate() call that produced the answer, from its prefill pass, so every vector belongs to the labelled answer. Layer ℓ is the residual stream after decoder block ℓ . We cache every layer (Qwen: 28 layers, $d = 3584$; LLaVA: 32 layers, $d = 4096$) at four pooled positions: **vision_mean** and **vision_max** (mean and elementwise max over the image tokens), **query_last** (the final prompt position, i.e. the state from which the first answer token is predicted) and **query_mean** (mean over the text tokens of the prompt). Pooling is what makes full-depth caching affordable: the image contributes on average 584 tokens per question for Qwen and 2,212 for LLaVA. Caches are stored in fp16 (1.5 GB for Qwen and 2.0 GB for LLaVA on the ChartQA test set).

4.3 Probes

A probe is an L_2 -regularised logistic regression on one pooled vector. For a chosen layer ℓ and position p , the input $h = h_\ell^p(x)$ is z-scored with the mean and standard deviation of the *training* items only, and the probe outputs a score $s(x) = w^\top h + b$, trained by minimising the logistic loss plus $\frac{\lambda}{2}\|w\|^2$ (L-BFGS; the intercept is not penalised). There are $d + 1$ parameters (3,585 for Qwen, 4,097 for LLaVA), which is more than the number of posi-

Model	Overall	Human	Augmented
Qwen2.5-VL-7B	87.3	79.8	94.8
LLaVA-NeXT-7B	52.9	44.6	61.1

Table 1: Relaxed accuracy (%) on ChartQA test.

tive training examples in every task, so λ is selected from $\{1, 10, 10^2, 10^3, 10^4\}$. Probes are scored by AUROC: the probability that a randomly chosen positive item receives a higher score than a randomly chosen negative one. We train two kinds:

Single-split sweep probes (binary, ChartQA structural, and absent-versus-present on synthetic data). Items are split 60/20/20 by *figure* (hashed, seed 42), so no chart appears in two splits. Every (position, layer, λ) cell is fitted on train and scored on validation; the best cell is chosen on validation only, and the test set is scored once.

Cross-validated type probes (the E3 matrix on synthetic data, where misreads are too few for a single held-out split). Items are assigned to five folds by figure. Within each outer fold, layer and λ are chosen by an inner cross-validation over the other four folds, and the probe then scores the held-out fold. Every item thus receives an out-of-fold score from a probe that never saw its chart. Position is fixed to `query_last` in advance.

All intervals are 95% percentile intervals from a bootstrap that resamples whole figures (2,000 draws), because questions on the same chart are correlated. Every probe is reported next to a baseline that sees no activations.

5 Naturalistic Arm: ChartQA

5.1 Accuracy

Table 1 gives relaxed accuracy on the ChartQA test set (1,250 human-written and 1,250 machine-augmented questions over 1,509 charts). Qwen is wrong on 317 questions and LLaVA on 1,178.

5.2 Labelling Qwen’s errors

The proposal planned two classes plus AMBIGUOUS. A pilot of 20 errors showed that this forced many items into AMBIGUOUS for reasons unrelated to the model, so the team adopted five labels: STRUCTURAL (the answer is explained by misreading the figure: wrong bar or series, misread digits, imprecise reading, miscounting), FABRICATION (content the figure does not support), COMPUTATION (inputs read correctly, arithmetic or logic wrong), NOT_AN_ERROR (the answer is

Label	n	Breakdown
NOT_AN_ERROR	112	51 format, 38 valid reading, 23 wrong gold
COMPUTATION	93	
STRUCTURAL	71	
AMBIGUOUS	40	14 question, 11 wrong gold, 12 other, 3 unreadable
FABRICATION	1	

Table 2: Labels for Qwen’s 317 ChartQA test errors.

actually right: an equivalent format, a wrong gold answer, or a valid reading of an ambiguous question) and AMBIGUOUS. Only STRUCTURAL and FABRICATION train type probes. All 317 errors were labelled by an LLM annotator (Claude) that saw only the chart image, question, gold answer and model answer, following the written rubric; the planned two-person human agreement check ($\kappa > 0.6$ on 200 items) has not yet been run.

Table 2 gives two findings. First, **fabrication is essentially absent** from Qwen’s ChartQA errors: one item in 317. ChartQA questions are written about content the chart contains, so a capable model rarely has occasion to invent content; its errors are misreads and arithmetic. The naturalistic arm therefore cannot supply fabrication examples, which is the risk the synthetic arm was planned for. Second, **35% of the scored errors are not errors**: most often the answer is correct in another format (“1:2” for 0.5), the gold answer is wrong, or the question admits the model’s reading. This matters for every probe trained on relaxed-accuracy labels.

What the error types look like. Examples from the labels, with the annotator’s rationale paraphrased:

- STRUCTURAL (71, 22%): a pie chart with seven labelled sectors counted as six; the adjacent 2017 bar reported instead of the asked 2018 bar; the top of a stacked segment reported instead of its height; a reversed comparison of two bars of similar length.
- COMPUTATION (93, 29%, the largest class of real errors): every input is printed and readable, but the operation is wrong, e.g. a ratio of levels (83/11) where the question asks for a ratio of changes, a sum over four printed values that matches no subset of years, or picking the larger of two printed changes wrongly.
- NOT_AN_ERROR (112, 35%): *format*: 0.95 for a gold of 95%, “increase” for “increasing”,

	Probe	Surface
Qwen, all errors	0.877 [.83,.92]	0.715
Qwen, without NOT_AN_ERROR	0.932 [.90,.96]	0.786
LLaVA, all errors	0.798 [.76,.84]	0.608
Qwen, structural vs. correct	0.905 [.84,.97]	–

Table 3: Test AUROC [95% CI] of the binary probe (wrong vs. right) and of the Qwen structural probe on ChartQA, against a surface-only baseline (question source, question length, number of image tokens, numeric gold). All selected query_last; Qwen layer 21 of 28, LLaVA layer 24 of 32.

or the same bar named by a different colour word; *gold wrong*: the model’s 28/19 = 1.474 is the correct ratio of the two bars while the gold 1.18 matches no ratio on the chart; *valid reading*: a question that names no year, answered for the latest year while the gold uses the first.

- AMBIGUOUS (40, 13%): both the gold and the model are wrong (11), the question has several readings and the model fits none (14), the needed value is unreadable, e.g. two near-identical shades of teal (3), or a misread cannot be told apart from an arithmetic slip (12, “other”).
- FABRICATION (1): pie slices of 35, 40, 13 and 11, where the answer 75 is a sum of two slices; the model’s 55 matches no slice and no sum.

Only 22% of Qwen’s scored errors are figure misreads; nearly as many are arithmetic or reasoning slips, which the two-class taxonomy of the proposal did not provide for.

5.3 Predicting whether the answer is wrong (E1)

Table 3 reports the binary probe. A sanity probe first confirmed that each cache decodes what it plainly contains (human-written vs. augmented question: validation AUROC 1.000 for Qwen, 0.999 for LLaVA). For both models the probe clearly exceeds the surface baseline, with non-overlapping intervals. Removing Qwen’s 112 NOT_AN_ERROR items raises the probe from 0.877 to 0.932: part of what looked like probe error was label noise. The probe is not reading question style, which alone predicts errors at 0.68 for Qwen because 80% of its errors are on human-written questions: within human-written questions only it reaches 0.895 (Qwen, cleaned) and 0.850 (LLaVA).




Figure 1: Binary probe, validation AUROC by relative depth (best λ per layer). Solid: last prompt token; dashed: mean over image tokens.

LLaVA’s lower raw AUROC should not be read as a weaker signal; it is wrong on 47% of questions against Qwen’s 13%, so the two probes separate different error sets at different base rates, and its gap over its own baseline (+0.19) is similar to Qwen’s (+0.15).

Which layers. Figure 1 shows validation AUROC by relative depth. In both models the signal is in the query-side state and builds with depth: query_last rises from about 0.70 at the first layer to a plateau over the last third of the network (Qwen 0.85 to 0.87 over layers 18 to 27, best at 21; LLaVA 0.83 to 0.84 over layers 20 to 28, best at 24). The image-token positions stay flat at 0.62 to 0.68 at every depth, close to the surface baseline. This differs from HALP’s finding that Qwen is best served by visual-only features for object hallucination on photographs; for chart errors the informative state is the one that has already read the question.

A structural probe on real charts. Using the 71 STRUCTURAL labels against all correct answers (other labels excluded), a Qwen probe reaches test AUROC 0.905, again at query_last, layer 21. With only 13 structural errors in the test split the interval [0.84, 0.97] is the result, and this probe does not yet have its own surface baseline. With one fabrication, the corresponding fabrication probe cannot be trained on ChartQA.

6 Synthetic Arm

6.1 Design

The synthetic arm makes fabrication happen on purpose, by asking about categories that are absent from the chart. Two design constraints come from earlier checks. (1) A probe comparing ChartQA

misreads with synthetic fabrications would learn “real vs. synthetic chart”. (2) On synthetic data the question template alone predicts the intended failure type at AUROC 1.000, so a probe comparing one template with another can read the template. Every contrast is therefore made *within one question template*, and every outcome is scored from the model’s actual answer, never from what the question was designed to induce.

Pilot 1 rendered 300 bar charts (matplotlib; 4 to 7 bars; six themes such as fruits or cities; values on 0–50 or 0–100 axes). Each chart is asked the same question twice, in identical words: once about a bar that is shown (“What is the value of Mango?”) and once about a plausible name from the same theme that is not (“... of Guava?”). **Pilot 2** (300 charts, 1,800 questions) added three changes after pilot 1. (i) Category names come in lookalike pairs (Dubai/Dublin, Eagle/Falcon, Lute/Flute); a chart shows at most one of a pair, and half the absent questions ask for the partner of a shown bar. (ii) Three question families, each asked about a shown and a missing name: *value*, *compare* (“Which is larger, Guava or Apple?”) and *neighbor* (“Which bar is immediately to the right of Guava?”). Naming any bar in reply to the last is a fabrication with no numeric escape. (iii) Harder charts: values never multiples of 5, axes up to 200, sparse ticks. Generation is deterministic, and both models were run on byte-identical charts.

Answers to missing names are scored as *rejected* (“None”), *fabricated* (any other number; naming the missing bar, its lookalike or any bar for neighbor questions), *zero* (“0”) or *picked-present* (naming the shown bar in a compare question). The team decided that “0” to a value question is a refusal expressed as a number, not a fabrication: no bar in these charts has zero height, so the model is not reporting something it read, and Qwen says “0” mostly when no bar resembles the asked name (141 of 157 unrelated names). It still counts as a wrong answer.

6.2 How the models answered

Table 4 summarises the answers. **Neither model refuses:** Qwen rejects 15 of 1,200 missing-name questions and LLaVA none. Qwen answers most missing values with “0”; when it does invent a number, 57 of its 76 pilot-2 inventions are the value of the lookalike bar, i.e. it binds the missing name to the nearest-looking bar that is present (asked for Dubai, it reports Dublin). LLaVA invents far

		Qwen	LLaVA
Shown bar, right	value (P1+P2)	554/600	185/600
	compare	300/300	209/300
	neighbor	294/300	255/300
Missing, value	rejected	15	0
	“0”	494	232
	fabricated	91	368
Missing, compare	fabricated	37	101
Missing, neighbor	fabricated	300	290

Table 4: Synthetic outcomes (pilots 1 and 2). Value rows pool both pilots (600 missing-name questions); compare and neighbor come from pilot 2 (300 each).

Family	Qwen	LLaVA
value	0.999 (L4)	0.977 (L12)
compare	0.991 (L4)	0.959 (L19)
neighbor	1.000 (L5)	0.996 (L18)
lookalike names only	0.980–1.000	0.922–0.991
text-only baseline	0.49–0.52	0.49–0.52
image-token probe	0.500	0.500

Table 5: Test AUROC of probes separating a missing from a shown name under identical wording (pilot 2; 62 charts per class in test). The image-token probe is a built-in control: image tokens precede the question, so they cannot see it.

more often (368 vs. 91), mostly round numbers near some bar. The models also differ on shown bars: Qwen reads 92% of values within tolerance, LLaVA 31%, because LLaVA snaps its readings to the nearest gridline (52 becomes 50). This yields four classes for the value template: C correct read, S misread (structural), F invented number (fabrication) and Z “0” (Qwen 554/46/91/494; LLaVA 185/415/368/232).

6.3 Do the models know the bar is missing?

Table 5 shows that **both models represent that the asked name is missing, although they answer anyway**. Two controls make the result interpretable. A classifier on the question text alone is at chance, because every name is a shown bar on some charts and missing on others. Probes on the image-token positions score exactly 0.500: the image precedes the question in the prompt and attention is causal, so both questions of a chart receive identical image vectors. The signal holds for lookalike names, and it appears early in Qwen (validation 0.99 from layer 2) and later in LLaVA (about layer 9). The absence probe does not respond to misreads of shown bars (Qwen 0.46, LLaVA 0.51 on shown-bar value questions).

6.4 Type probes and cross-transfer (E3)

Table 6 is the central result. Each probe is trained on one pair of classes and then scored on every pair. If misreads and fabrications leave different internal signatures, each probe should be strong on its own task (bold) and weak on the other failure type.

Qwen: largely type-specific. The structural probe separates misreads from correct reads (0.887) but barely responds to fabrications (0.594); the fabrication probe is perfect on its own task and near chance on misreads (0.576, interval touching 0.5). The structural probe exceeds its metadata baseline by a modest 0.07 with overlapping intervals: much of Qwen’s misread risk is chart difficulty visible in the settings (0–200 axes with sparse ticks account for most misreads).

LLaVA: the misread signal is not type-specific. Its structural probe fires *more* on fabrications (0.879) and on “0” answers (0.987) than on its own task (0.776). It behaves like a general “this answer is not grounded” direction, strongest when there is no bar at all. Moreover, because LLaVA snaps to gridlines, the distance of the true value from the nearest multiple of 10 alone predicts its misreads at 0.852, above the probe; for Qwen the same number gives only 0.616, far below its probe. (This feature was chosen after seeing LLaVA’s rounding, so it is a diagnostic rather than a pre-registered baseline.)

In both models the “fabrication” probe is a missing-bar detector. It scores “0” answers exactly as high as invented numbers (Z vs C 1.000) and even ranks them above (F vs Z below 0.3). What separates an invented number from “0” is decodable (0.984, 0.986), but only in the last layers (Qwen 22 to 27, LLaVA 23 to 30), where it largely reads out the answer about to be produced rather than giving an early warning.

Which layers. Figure 2 shows own-task AUROC by relative depth. The missing-bar signal saturates early (Qwen by layer 3, LLaVA by about layer 8). The misread signal builds gradually and peaks in the last third (layers chosen by inner CV: Qwen 16 to 25, LLaVA 20 to 30).

7 Discussion

Is the hypothesis supported? Partly. For Qwen2.5-VL-7B, the state before a misread and the state before a missing-bar fabrication are linearly distinguishable, and each probe is largely blind to the other failure type. For LLaVA-NeXT-7B only one direction holds: the missing-bar signal is spe-

Probe (trained on)	Qwen2.5-VL-7B				LLaVA-NeXT-7B			
	S vs C	F vs C	Z vs C	F vs Z	S vs C	F vs C	Z vs C	F vs Z
structural (S vs C)	0.887	0.594	0.673	0.400	0.776	0.879	0.987	0.215
metadata baseline	0.814	0.557	0.515	0.543	0.512	0.522	0.486	0.537
fabrication (F vs C)	0.576	1.000	1.000	0.286	0.680	0.988	1.000	0.243
metadata baseline	0.588	0.840	0.564	0.787	0.642	0.769	0.643	0.651
invented vs. zero (F vs Z)	0.336	0.138	0.000	0.984	0.520	0.226	0.000	0.986

Table 6: Cross-transfer matrix on synthetic value questions (5-fold CV by chart, out-of-fold AUROC). Rows: the probe and what it was trained on; columns: the contrast it is scored on. Bold: each probe on its own task. Metadata baseline: the same procedure on chart settings only (pilot, axis range, tick density, bar count, lookalike present, phrasing). 95% CIs for the bold cells: Qwen [0.82, 0.94], [1.00, 1.00], [0.97, 0.99]; LLaVA [0.73, 0.82], [0.98, 0.99], [0.98, 0.99]; for the structural probe on F vs C, Qwen [0.55, 0.65], LLaVA [0.85, 0.90].




Figure 2: Type probes on synthetic value questions: out-of-fold AUROC on their own task by relative depth (λ chosen per fold by inner CV).

cific, but the misread probe captures a general lack of grounding that is strongest on fabrications, and a chart-level feature explains its misreads better. Under the criterion we set in the proposal (“strong cross-transfer indicates one signal wearing two labels”), LLaVA shows exactly that, from the structural side.

What the fabrication signal is. With our data, a fabrication is defined by the missing category, so “predicting fabrication” and “detecting that the asked category is missing” cannot be separated. The finding we can state is that both models encode, early and linearly, that the question is about something not on the chart, and then answer it anyway. This is the signal a controller needs: it identifies exactly the questions where re-reading cannot help and abstention is the right repair.

How the models differ. Qwen is accurate (87% on ChartQA, 92% on synthetic reads), rarely invents numbers, and when it does it binds the missing name to a lookalike bar. LLaVA is less accurate (53% and 31%), reads coarsely to gridlines, and

invents numbers four times as often. Qwen forms the missing-bar representation within its first few layers; LLaVA forms it in the first third. Both concentrate error information in the query-side state of the last third of the network rather than in the image tokens.

Validity. Every number above sits next to a baseline with no access to activations: a surface baseline on ChartQA, a text-only baseline and an image-token control on synthetic data, and a metadata baseline for the type probes. The image-token control is exact by construction, which rules out leakage through the image. Remaining threats are listed in Section 9.

8 Progress Against Plan and Timeline

Done. Inference and full-depth caching for both models on ChartQA test (2,500), ChartQA validation (960, Qwen) and both synthetic pilots (2,400); Qwen error labels (317); E1 for both models with surface baselines; the E3 matrix for both models on synthetic data; the within-source and image-token controls.

Changed. The LLM judge was replaced by an LLM annotator with a five-label rubric (adding COMPUTATION and NOT_AN_ERROR); synthetic fabrication is induced with missing categories and lookalike names, while the proposal’s out-of-range quantities are not yet implemented; the planned fabricated-versus-rejected contrast was replaced by missing-versus-shown, since neither model rejects often enough to train it; ChartGemma is dropped.

Go/no-go (30 Sep). The probes separate the types for Qwen and the missing-bar signal survives every control in both models, so the second half pursues E4 and E5 rather than a purely negative write-up.

Remaining.

- **Oct 1–8:** human agreement check (κ on 200 items, two annotators); labels for LLaVA’s 1,178 ChartQA errors; residualised probes (E2) and a surface baseline for the ChartQA structural probe.
- **Oct 1–15:** E4 cross-intervention matrix (question-conditioned crop, verify prompt, abstain) on synthetic items, whose bar geometry is recorded at render time, for both models.
- **Oct 16–25:** E5 controller from the binary and missing-bar probes; risk–coverage and AURC against binary selective prediction and the oracle.
- **Oct 26–31:** final report.

9 Limitations

The ChartQA labels come from a single LLM annotator and have not yet passed the human agreement check. Fabrication is absent from ChartQA, so the type contrast rests on synthetic bar charts and on one question template, where fabrication coincides with a missing category. Qwen has only 46 synthetic misreads, and the two models’ type matrices rest on very different class balances. The ChartQA structural probe has 13 positives in test and no surface baseline yet. Probes are linear and diagnostic; whether the signals support a useful intervention is the question E4 and E5 must answer.

References

- Shir Ashury-Tahan et al. 2026. ErrorMap and ErrorAtlas: Charting the failure landscape of large language models. *arXiv preprint arXiv:2601.15812*.
- Shuai Bai et al. 2025. Qwen2.5-VL technical report. *arXiv preprint arXiv:2502.13923*.
- Anirudh Iyengar Kaniyar Narayana Iyengar, Tampu Ravi Kumar, Gaurav Najpande, Manan Suri, Dinesh Manocha, Puneet Mathur, and Vivek Gupta. 2026. DRAGON: A benchmark for evidence-grounded visual reasoning over diagrams. *arXiv preprint arXiv:2604.25231*.
- Sai Akhil Kogilathota, Sripadha Vallabha E G, Luzhe Sun, and Jiawei Zhou. 2026. HALP: Detecting hallucinations in vision-language models without generating a single token. In *Proceedings of the 19th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 6067–6085.
- Jinsong Li, Xiaoyi Dong, Yuhang Zang, Yuhang Cao, Jiaqi Wang, and Dahua Lin. 2026. Visual self-refine: A pixel-guided paradigm for accurate chart parsing. In *International Conference on Learning Representations (ICLR)*.
- Haotian Liu, Chunyuan Li, Yuheng Li, Bo Li, Yuanhan Zhang, Sheng Shen, and Yong Jae Lee. 2024. [LLaVA-NeXT: Improved reasoning, OCR, and world knowledge](#).
- Jiaxin Liu, Ding Zhong, Yue Wang, Zhidong Yang, Zhaolu Kang, Guangyuan Dong, Qishi Zhan, Pengcheng Fang, and Aofan Liu. 2026. Dual-pathway circuits of object hallucination in vision-language models. *arXiv preprint arXiv:2605.13156*.
- Ahmed Masry, Do Xuan Long, Jia Qing Tan, Shafiq Joty, and Enamul Hoque. 2022. ChartQA: A benchmark for question answering about charts with visual and logical reasoning. In *Findings of the Association for Computational Linguistics: ACL 2022*.
- Anku Rani, Aparna Garimella, Apoorv Saxena, Balaji Vasanth Srinivasan, and Paul Pu Liang. 2025. RADAR: A reasoning-guided attribution framework for explainable visual data analysis. *arXiv preprint arXiv:2508.16850*.
- Subramanyam Sahoo, Vinija Jain, Aman Chadha, and Divya Chaudhary. 2026. Linear probes detect task format, not reasoning mode in language model hidden states. *arXiv preprint arXiv:2606.02907*.
- Aojie Yuan, Zhiyuan Julian Su, Haiyue Zhang, Yi Nian, and Yue Zhao. 2026. Hidden error awareness in chain-of-thought reasoning: The signal is diagnostic, not causal. *arXiv preprint arXiv:2605.09502*.
- Feiran Zhang, Yixin Wu, Zhenghua Wang, Xiaohua Wang, Changze Lv, Xuanjing Huang, and Xiaoqing Zheng. 2026. VIB-probe: Detecting and mitigating hallucinations in vision-language models via variational information bottleneck. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 23509–23521.